


R Exam

> #Q1. Write an R program to create a sequence of numbers from 20 to 50 and find the mean of numbers from 20 to 60 and the sum of numbers from 51 to 91.

```
> seq(20, 50)
[1] 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49
50
> mean(20:60)
[1] 40
> mean(seq(20, 60))
[1] 40
> sum(51:91)
[1] 2911
> sum(seq(51, 91))
[1] 2911
>
```

> #Q2. A student scored 70 marks in English, 95 marks in Science, 80 marks in Maths and 74 marks in History. Write an R program to plot a simple bar chart displaying the scores of the given subjects.

```
> subjects <- c("English"=70, "Science"=95, "Maths"=80, "History"=74)
> barplot(subjects, col=c("aquamarine", "green", "red", "orange"))
```



>

> #Q3. Write a R program to create a data frame to store the following details of 5 employees.

```
> name <- c("Anastasia S", "Dima R", "Katherine S", "JAMES A", "LAURA MARTIN")
> gender <- c("M", "M", "F", "F", "M")
> age <- c(23, 22, 25, 26, 32)
> desig <- c("Clerk", "Manager", "Executive", "CEO", "ASSISTANT")
> ssn <- c("123-34-2346", "23-44-779", "556-24-433", "123-98-987", "679-77-576")
> employees <- data.frame(name, gender, age, desig, ssn)
> colnames(employees) <- c("Name", "Gender", "Age", "Designation", "SSN")
> employees
```

	Name	Gender	Age	Designation	SSN
1	Anastasia S	M	23	Clerk	123-34-2346
2	Dima R	M	22	Manager	23-44-779
3	Katherine S	F	25	Executive	556-24-433
4	JAMES A	F	26	CEO	123-98-987
5	LAURA MARTIN	M	32	ASSISTANT	679-77-576

```
>
```

> #Q4. Write an R program to create a list of heterogeneous data, which includes character, numeric and logical vectors. Print the list.

```
> l <- list(c("Male", "Female"), c(24, 25), TRUE)
> l
[[1]]
[1] "Male" "Female"
```

```
[[2]]
[1] 24 25
```

```
[[3]]
[1] TRUE
```

```
>
```

> #Q.5 Write an R program to convert a given matrix to a 1-dimensional array.

```
> mat <- matrix(1:12, ncol=4)
> mat
     [,1] [,2] [,3] [,4]
[1,]  1   4   7  10
[2,]  2   5   8  11
[3,]  3   6   9  12
> array(as.vector(mat))
[1] 1 2 3 4 5 6 7 8 9 10 11 12
>
```

> #Q.6 Write a R program to create a list containing a given vector, a matrix, and a list and add an element at the end of the list

```
> li <- list(c("Red", "Green", "Black"), matrix(seq(1, 11, 2), ncol=3), list("Python", "PHP", "Java"))
```

```
> li
```

```
[[1]]
```

```
[1] "Red" "Green" "Black"
```

```
[[2]]
```

```
 [,1] [,2] [,3]
```

```
[1,]  1  5  9
```

```
[2,]  3  7 11
```

```
[[3]]
```

```
[[3]][[1]]
```

```
[1] "Python"
```

```
[[3]][[2]]
```

```
[1] "PHP"
```

```
[[3]][[3]]
```

```
[1] "Java"
```

```
> li <- append(li, 4)
```

```
> li
```

```
[[1]]
```

```
[1] "Red" "Green" "Black"
```

```
[[2]]
```

```
 [,1] [,2] [,3]
```

```
[1,]  1  5  9
```

```
[2,]  3  7 11
```

```
[[3]]
```

```
[[3]][[1]]
```

```
[1] "Python"
```

```
[[3]][[2]]
```

```
[1] "PHP"
```

```
[[3]][[3]]
```

```
[1] "Java"
```

```
[[4]]
```

```
[1] 4
```

```
>
```

```
> #Q.7 Write an R program to merge two given lists into one list.
```

```
> List1= list(1, 2, 3)
```

```
> List2 = list("Red", "Green", "Black")
```

```
> c(List1, List2)
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] 2
```

```
[[3]]
```

```
[1] 3
```

```
[[4]]
```

```
[1] "Red"
```

```
[[5]]
```

```
[1] "Green"
```

```
[[6]]
```

```
[1] "Black"
```

```
>
```

```
> #Q.8 Write an R program to convert a given data frame to a list by rows.
```

```
> name <- c("Anastasia", "Dima", "Katherine", "James", "Emily", "Michael", "Matthew", "Laura",  
"Kevin")
```

```
> score <- c(12.5, 9.0, 16.5, 12.0, 9.0, 20.0, 14.5, 13.5, 8.0)
```

```
> attempts <- c(1, 3, 2, 3, 2, 3, 1, 1, 2)
```

```
> qualify <- c("yes", "no", "yes", "no", "no", "yes", "yes", "no", "no")
```

```
> df <- data.frame(name, score, attempts, qualify)
```

```
> colnames(df) <- c("Name", "Score", "attempts", "qualify")
```

```
> df
```

```
   Name Score attempts qualify
```

```
1 Anastasia 12.5     1    yes
```

```
2   Dima  9.0    3   no
3 Katherine 16.5    2  yes
4   James 12.0    3   no
5   Emily  9.0    2   no
6 Michael 20.0    3  yes
7 Matthew 14.5    1  yes
8   Laura 13.5    1   no
9   Kevin  8.0    2   no
```

```
> li2 <- list()
```

```
> li2 <- apply(df, 1, function(x) append(li2, x))
```

```
> li2
```

```
[[1]]
```

```
[[1]]$Name
```

```
[1] "Anastasia"
```

```
[[1]]$Score
```

```
[1] "12.5"
```

```
[[1]]$attempts
```

```
[1] "1"
```

```
[[1]]$qualify
```

```
[1] "yes"
```

```
[[2]]
```

```
[[2]]$Name
```

```
[1] "Dima"
```

```
[[2]]$Score
```

```
[1] "9.0"
```

```
[[2]]$attempts
```

```
[1] "3"
```

```
[[2]]$qualify
```

```
[1] "no"
```

```
[[3]]
```

```
[[3]]$Name
```

[1] "Katherine"

[[3]]\$Score

[1] "16.5"

[[3]]\$Attempts

[1] "2"

[[3]]\$qualify

[1] "yes"

[[4]]

[[4]]\$Name

[1] "James"

[[4]]\$Score

[1] "12.0"

[[4]]\$Attempts

[1] "3"

[[4]]\$qualify

[1] "no"

[[5]]

[[5]]\$Name

[1] "Emily"

[[5]]\$Score

[1] " 9.0"

[[5]]\$Attempts

[1] "2"

[[5]]\$qualify

[1] "no"

[[6]]

[[6]]\$Name
[1] "Michael"

[[6]]\$Score
[1] "20.0"

[[6]]\$Attempts
[1] "3"

[[6]]\$qualify
[1] "yes"

[[7]]
[[7]]\$Name
[1] "Matthew"

[[7]]\$Score
[1] "14.5"

[[7]]\$Attempts
[1] "1"

[[7]]\$qualify
[1] "yes"

[[8]]
[[8]]\$Name
[1] "Laura"

[[8]]\$Score
[1] "13.5"

[[8]]\$Attempts
[1] "1"

[[8]]\$qualify
[1] "no"

```
[[9]]
[[9]]$Name
[1] "Kevin"
```

```
[[9]]$Score
[1] " 8.0"
```

```
[[9]]$attempts
[1] "2"
```

```
[[9]]$qualify
[1] "no"
```

```
>
```

> #Q.9 Write an R program to create a correlation matrix from a data frame of the same data type.

```
> d <- data.frame(x1=rnorm(5), x2=rnorm(5), x3=rnorm(5))
> d
```

```
      x1      x2      x3
1 -0.89691455 0.1324203 0.4176508
2 0.18484918 0.7079547 0.9817528
3 1.58784533 -0.2396980 -0.3926954
4 -1.13037567 1.9844739 -1.0396690
5 -0.08025176 -0.1387870 1.7822290
```

```
> cor(d)
```

```
      x1      x2      x3
x1 1.0000000 -0.6205198 0.0696402
x2 -0.6205198 1.0000000 -0.5720479
x3 0.0696402 -0.5720479 1.0000000
```

```
>
```

> #Q.10 Write an R program to rotate a given matrix 90 degrees clockwise.

```
> mt <- matrix(1:9, ncol=3)
```

```
> mt
```

```
      [,1] [,2] [,3]
[1,]  1  4  7
[2,]  2  5  8
[3,]  3  6  9
```

```
> clockwise <- function(x){ t(apply(x, 2, rev)) }
```

```
> clockwise(clockwise(mt))
```

```
      [,1] [,2] [,3]
```



```
[1,] 9 6 3
[2,] 8 5 2
[3,] 7 4 1
```

```
>
```

> #Q.11 Check for missing values in the 'mtcars' data set.

```
> sum(is.na(mtcars))
```

```
[1] 0
```

```
>
```

> #Q.12 Check which attributes are important to determine the mpg of a car in the 'mtcars' data set.

```
> str(mtcars)
```

```
'data.frame': 32 obs. of 11 variables:
```

```
$ mpg : num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
```

```
$ cyl : num 6 6 4 6 8 6 8 4 4 6 ...
```

```
$ disp: num 160 160 108 258 360 ...
```

```
$ hp : num 110 110 93 110 175 105 245 62 95 123 ...
```

```
$ drat: num 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
```

```
$ wt : num 2.62 2.88 2.32 3.21 3.44 ...
```

```
$ qsec: num 16.5 17 18.6 19.4 17 ...
```

```
$ vs : num 0 0 1 1 0 1 0 1 1 1 ...
```

```
$ am : num 1 1 1 0 0 0 0 0 0 0 ...
```

```
$ gear: num 4 4 4 3 3 3 3 4 4 4 ...
```

```
$ carb: num 4 4 1 1 2 1 4 2 2 4 ...
```

```
> library(caret)
```

```
Loading required package: ggplot2
```

```
Need help getting started? Try the R Graphics Cookbook: https://r-graphics.org
```

```
Loading required package: lattice
```

```
> reg_model1 <- lm(mpg ~ ., data=mtcars)
```

```
> summary(reg_model1)
```

Call:

```
lm(formula = mpg ~ ., data = mtcars)
```

Residuals:

```
    Min     1Q  Median     3Q     Max
-3.4506 -1.6044 -0.1196  1.2193  4.6271
```

Coefficients:

```
      Estimate Std. Error t value Pr(>|t|)
(Intercept) 12.30337   18.71788   0.657  0.5181
cyl         -0.11144    1.04502  -0.107  0.9161
```

disp	0.01334	0.01786	0.747	0.4635
hp	-0.02148	0.02177	-0.987	0.3350
drat	0.78711	1.63537	0.481	0.6353
wt	-3.71530	1.89441	-1.961	0.0633 .
qsec	0.82104	0.73084	1.123	0.2739
vs	0.31776	2.10451	0.151	0.8814
am	2.52023	2.05665	1.225	0.2340
gear	0.65541	1.49326	0.439	0.6652
carb	-0.19942	0.82875	-0.241	0.8122

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.65 on 21 degrees of freedom

Multiple R-squared: 0.869, Adjusted R-squared: 0.8066

F-statistic: 13.93 on 10 and 21 DF, p-value: 3.793e-07

```
> reg_model2 <- lm(mpg ~ cyl+hp+drat+wt+qsec+vs+am+gear+carb, data=mtcars)
```

```
> summary(reg_model2)
```

Call:

```
lm(formula = mpg ~ cyl + hp + drat + wt + qsec + vs + am + gear +
    carb, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.7863	-1.4055	-0.2635	1.2029	4.4753

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.55052	18.52585	0.677	0.5052
cyl	0.09627	0.99715	0.097	0.9240
hp	-0.01295	0.01834	-0.706	0.4876
drat	0.92864	1.60794	0.578	0.5694
wt	-2.62694	1.19800	-2.193	0.0392 *
qsec	0.66523	0.69335	0.959	0.3478
vs	0.16035	2.07277	0.077	0.9390
am	2.47882	2.03513	1.218	0.2361
gear	0.74300	1.47360	0.504	0.6191
carb	-0.61686	0.60566	-1.018	0.3195

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.623 on 22 degrees of freedom
Multiple R-squared: 0.8655, Adjusted R-squared: 0.8105
F-statistic: 15.73 on 9 and 22 DF, p-value: 1.183e-07

```
> reg_model3 <- lm(mpg ~ cyl+wt+qsec+wt+am+gear+carb, data=mtcars)
> summary(reg_model3)
```

Call:

```
lm(formula = mpg ~ cyl + wt + qsec + wt + am + gear + carb, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.2148	-1.1992	-0.2412	1.4018	4.4595

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	16.2624	15.9126	1.022	0.3166
cyl	-0.3137	0.7971	-0.393	0.6973
wt	-2.9548	1.0611	-2.785	0.0101 *
qsec	0.7695	0.5853	1.315	0.2005
am	2.6522	1.8807	1.410	0.1708
gear	0.6415	1.3835	0.464	0.6469
carb	-0.6764	0.5481	-1.234	0.2287

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.509 on 25 degrees of freedom
Multiple R-squared: 0.8602, Adjusted R-squared: 0.8267
F-statistic: 25.64 on 6 and 25 DF, p-value: 1.542e-09

```
> reg_model4 <- lm(mpg ~ cyl+qsec+am+wt+carb, data=mtcars)
> summary(reg_model4)
```

Call:

```
lm(formula = mpg ~ cyl + qsec + am + wt + carb, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.2795	-1.2098	-0.3826	1.3961	4.4050

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	20.0379	13.4638	1.488	0.14871
cyl	-0.4612	0.7198	-0.641	0.52728
qsec	0.7272	0.5693	1.277	0.21276
am	2.9417	1.7471	1.684	0.10419
wt	-3.0462	1.0268	-2.967	0.00638 **
carb	-0.5222	0.4291	-1.217	0.23456

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.471 on 26 degrees of freedom

Multiple R-squared: 0.859, Adjusted R-squared: 0.8319

F-statistic: 31.69 on 5 and 26 DF, p-value: 2.847e-10

```
> reg_model5 <- lm(mpg ~ qsec+gear+wt+carb, data=mtcars)
```

```
> summary(reg_model5)
```

Call:

```
lm(formula = mpg ~ qsec + gear + wt + carb, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.5079	-1.9616	-0.0451	0.8937	5.2861

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	13.5291	8.1655	1.657	0.109125
qsec	0.7613	0.3462	2.199	0.036626 *
gear	1.9228	1.0973	1.752	0.091081 .
wt	-3.7019	0.8902	-4.159	0.000291 ***
carb	-0.7848	0.5470	-1.435	0.162841

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.543 on 27 degrees of freedom

Multiple R-squared: 0.8449, Adjusted R-squared: 0.8219

F-statistic: 36.77 on 4 and 27 DF, p-value: 1.471e-10

```
> reg_model6 <- lm(mpg ~ qsec+gear+wt, data=mtcars)
```

```
> summary(reg_model6)
```

Call:

```
lm(formula = mpg ~ qsec + gear + wt, data = mtcars)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.7178	-1.7915	-0.3533	1.1897	5.6333

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.9432	8.3080	1.558	0.13049
qsec	1.0486	0.2877	3.645	0.00108 **
gear	0.8914	0.8446	1.055	0.30025
wt	-4.6178	0.6320	-7.306	5.91e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.591 on 28 degrees of freedom

Multiple R-squared: 0.8331, Adjusted R-squared: 0.8152

F-statistic: 46.57 on 3 and 28 DF, p-value: 5.2e-11

> #After comparing the Adjusted R squared, Error and R squared value selecting model 5

>

> **#Q.13 Build a simple linear model to predict the mpg of a car in the 'mtcars' data set.**

```
> final_reg_model <- lm(mpg ~ qsec+gear+wt+carb, data=mtcars)
```

```
> prediction <- predict(final_reg_model, newdata=data.frame(qsec=16.5, wt=2.62, gear=4, carb=4))
```

```
> prediction
```

1

20.94371

>

> **#Q.14 Build a logistic regression model using the glm function to know the effect of admission into graduate school. The target variable,**

> **#admit/don't admit, is a binary variable Use the given "binary.csv" dataset.**

```
> setwd("C:/zubeda/PGA02_Zubu/R Programming/R Exam/Dataset")
```

```
> admission <- read.csv("binary.csv")
```

```
> head(admission)
```

	admit	gre	gpa	rank
--	-------	-----	-----	------

1	0	380	3.61	3
---	---	-----	------	---

2	1	660	3.67	3
---	---	-----	------	---

3	1	800	4.00	1
---	---	-----	------	---

```

4  1 640 3.19  4
5  0 520 2.93  4
6  1 760 3.00  2
> str(admission)
'data.frame':  400 obs. of  4 variables:
 $ admit: int  0 1 1 1 0 1 1 0 1 0 ...
 $ gre  : int  380 660 800 640 520 760 560 400 540 700 ...
 $ gpa  : num  3.61 3.67 4 3.19 2.93 3 2.98 3.08 3.39 3.92 ...
 $ rank : int  3 3 1 4 4 2 1 2 3 2 ...
> admission$rank <- as.factor(admission$rank)
> logit <- glm(admit ~ gre+gpa+rank, data=admission, family="binomial")
> summary(logit)

```

Call:

```

glm(formula = admit ~ gre + gpa + rank, family = "binomial",
    data = admission)

```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.6268	-0.8662	-0.6388	1.1490	2.0790

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.989979	1.139951	-3.500	0.000465 ***
gre	0.002264	0.001094	2.070	0.038465 *
gpa	0.804038	0.331819	2.423	0.015388 *
rank2	-0.675443	0.316490	-2.134	0.032829 *
rank3	-1.340204	0.345306	-3.881	0.000104 ***
rank4	-1.551464	0.417832	-3.713	0.000205 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 499.98 on 399 degrees of freedom
Residual deviance: 458.52 on 394 degrees of freedom
AIC: 470.52

Number of Fisher Scoring iterations: 4

>

> #Q.15 Use the given variables from the titanic dataset and build the decision tree on train data. Variables from dataset: survived, embarked, sex, sibsp, parch, fare

```
> titanic <- read.csv("Titanic_data.csv")
```

```
> head(titanic)
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
1	0	3	Braund, Mr. Owen Harris	male	22	1	0	A/5 21171	7.2500	
2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38	1	0	PC 17599	71.2833	C85
3	0	3	Heikkinen, Miss. Laina	female	26	0	0	STON/O2.	7.9250	3101282
4	0	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35	1	0		53.1000	C123
5	0	3	Allen, Mr. William Henry	male	35	0	0		8.0500	373450
6	0	3	Moran, Mr. James	male	NA	0	0		8.4583	330877

```
Embarked
```

```
1 S
2 C
3 S
4 S
5 S
6 Q
```

```
> titanic_df <- titanic[, c("Survived", "Embarked", "Sex", "SibSp", "Parch", "Fare")]
```

```
> head(titanic_df)
```

	Survived	Embarked	Sex	SibSp	Parch	Fare
1	0	S	male	1	0	7.2500
2	1	C	female	1	0	71.2833
3	0	S	female	0	0	7.9250
4	0	S	female	1	0	53.1000
5	0	S	male	0	0	8.0500
6	0	Q	male	0	0	8.4583

```
> library(rpart)
```

```
> library(caTools)
```

```
> set.seed(123)
```

```
> #Splitting dataset
```

```
> split <- sample.split(titanic_df, SplitRatio=0.8)
```

```
> training <- subset(titanic_df, split == TRUE)
```

```
> test <- subset(titanic_df, split == FALSE)
```

```

> dim(training)
[1] 595  6
> dim(test)
[1] 296  6
> training$Survived <- as.factor(training$Survived)
> test$Survived <- as.factor(test$Survived)
> #Building Decision Tree model
> mtree <- rpart(Survived ~ ., data=training, method="class")
> mtree
n= 595

```

node), split, n, loss, yval, (yprob)

* denotes terminal node

```

1) root 595 214 0 (0.64033613 0.35966387)
 2) Sex=male 393  72 0 (0.81679389 0.18320611) *
 3) Sex=female 202  60 1 (0.29702970 0.70297030)
 6) SibSp>=2.5 18  5 0 (0.72222222 0.27777778) *
 7) SibSp< 2.5 184  47 1 (0.25543478 0.74456522)
14) Fare< 25.075 105  41 1 (0.39047619 0.60952381)
28) Embarked=S 62  29 1 (0.46774194 0.53225806)
 56) Fare< 10.48125 25  8 0 (0.68000000 0.32000000) *
 57) Fare>=10.48125 37  12 1 (0.32432432 0.67567568) *
29) Embarked=C,Q 43  12 1 (0.27906977 0.72093023)
 58) Fare< 15.3729 33  12 1 (0.36363636 0.63636364)
116) Fare>=14.15625 8  1 0 (0.87500000 0.12500000) *
117) Fare< 14.15625 25  5 1 (0.20000000 0.80000000) *
 59) Fare>=15.3729 10  0 1 (0.00000000 1.00000000) *
15) Fare>=25.075 79  6 1 (0.07594937 0.92405063) *

```

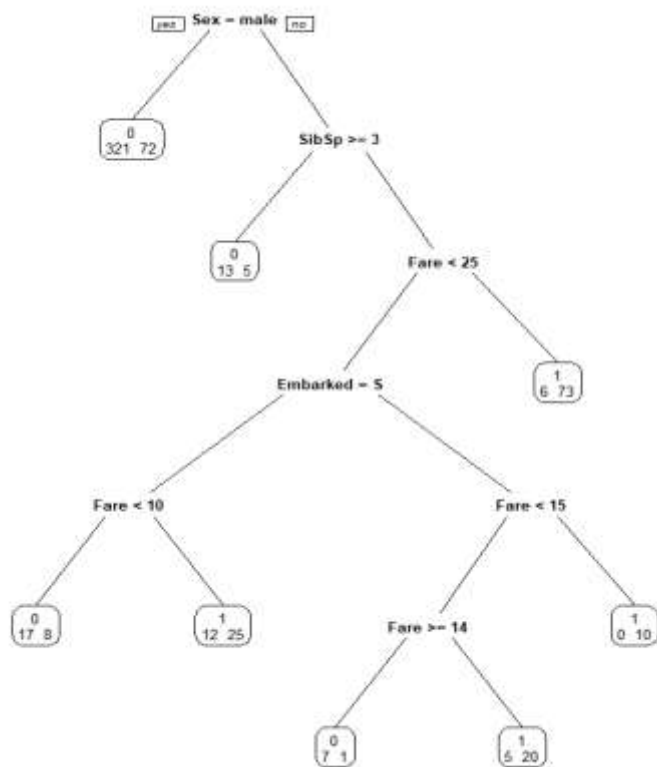
>

> #Q.16 Create a plot to display the result of decision tree.

```

> library(rpart.plot)
> prp(mtree, faclen=0, extra=1, cex=0.8)

```

>

> #Q.17 Create the confusion matrix for the above model.

> prediction <- predict(mtree, newdata=test[, -1], type="class")

> prediction

```

 4  5 10 11 16 17 22 23 28 29 34 35 40 41 46 47 52 53 58 59 64 65 70 71 76 77 82
83 88 89 94 95 100 101 106 107
 1  0  1  1  1  0  0  1  0  1  0  0  1  0  0  0  0  1  0  1  0  0  0  0  0  0  1  0  0  0  0  0
0  0  0
112 113 118 119 124 125 130 131 136 137 142 143 148 149 154 155 160 161 166 167 172 173
178 179 184 185 190 191 196 197 202 203 208 209 214 215
 0  0  0  0  1  0  0  0  0  1  0  1  1  0  0  0  0  0  1  0  1  1  0  0  1  0  1  1  0  0  0  0
1  0  0
220 221 226 227 232 233 238 239 244 245 250 251 256 257 262 263 268 269 274 275 280 281
286 287 292 293 298 299 304 305 310 311 316 317 322 323
 0  0  0  0  0  0  1  0  0  0  0  0  0  1  0  0  0  1  0  1  1  0  0  0  1  0  1  0  1  1  0
1  0  1
328 329 334 335 340 341 346 347 352 353 358 359 364 365 370 371 376 377 382 383 388 389
394 395 400 401 406 407 412 413 418 419 424 425 430 431
 1  1  0  1  0  0  1  1  0  0  1  1  0  0  1  0  1  0  1  0  1  1  1  0  0  0  0  1  1  0  1
0  0  0

```

```

436 437 442 443 448 449 454 455 460 461 466 467 472 473 478 479 484 485 490 491 496 497
502 503 508 509 514 515 520 521 526 527 532 533 538 539
  1  1  0  0  0  1  0  0  0  0  0  0  0  1  0  0  0  0  0  0  1  1  1  0  0  1  0  0  1  0  1  0
0  1  0
544 545 550 551 556 557 562 563 568 569 574 575 580 581 586 587 592 593 598 599 604 605
610 611 616 617 622 623 628 629 634 635 640 641 646 647
  0  0  0  0  0  1  0  0  1  0  1  0  0  1  1  0  1  0  0  0  0  0  1  1  1  0  0  0  1  0  0  0  0
0  0  0
652 653 658 659 664 665 670 671 676 677 682 683 688 689 694 695 700 701 706 707 712 713
718 719 724 725 730 731 736 737 742 743 748 749 754 755
  1  0  1  0  0  0  1  1  0  0  0  0  0  0  0  0  1  0  1  0  0  1  0  0  0  0  1  0  1  0  1  1
0  0  1
760 761 766 767 772 773 778 779 784 785 790 791 796 797 802 803 808 809 814 815 820 821
826 827 832 833 838 839 844 845 850 851 856 857 862 863
  1  0  1  0  0  1  1  0  0  0  0  0  0  1  1  0  0  0  0  0  0  1  0  0  0  0  0  0  0  0  1  0  0
1  0  1
868 869 874 875 880 881 886 887
  0  0  0  1  1  1  1  0

```

Levels: 0 1

```
> confusionMatrix(data=prediction, reference=test$Survived, positive="1")
```

Confusion Matrix and Statistics

```

      Reference
Prediction 0  1
      0 154  45
      1  14  83

```

Accuracy : 0.8007

95% CI : (0.7506, 0.8447)

No Information Rate : 0.5676

P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.5819

Mcnemar's Test P-Value : 9.397e-05

Sensitivity : 0.6484

Specificity : 0.9167

Pos Pred Value : 0.8557

Neg Pred Value : 0.7739

Prevalence : 0.4324

Detection Rate : 0.2804
Detection Prevalence : 0.3277
Balanced Accuracy : 0.7826

'Positive' Class : 1

>

> #Q.18 Perform k-means clustering on USArrest dataset. Scale the data before performing clustering.

> arrests <- USArrests

> head(arrests)

	Murder	Assault	UrbanPop	Rape
Alabama	13.2	236	58	21.2
Alaska	10.0	263	48	44.5
Arizona	8.1	294	80	31.0
Arkansas	8.8	190	50	19.5
California	9.0	276	91	40.6
Colorado	7.9	204	78	38.7

> dim(arrests)

[1] 50 4

> str(arrests)

'data.frame': 50 obs. of 4 variables:

\$ Murder : num 13.2 10 8.1 8.8 9 7.9 3.3 5.9 15.4 17.4 ...

\$ Assault : int 236 263 294 190 276 204 110 238 335 211 ...

\$ UrbanPop: int 58 48 80 50 91 78 77 72 80 60 ...

\$ Rape : num 21.2 44.5 31 19.5 40.6 38.7 11.1 15.8 31.9 25.8 ...

> data_standardized <- data.frame(scale(arrests))

> head(data_standardized)

	Murder	Assault	UrbanPop	Rape
Alabama	1.24256408	0.7828393	-0.5209066	-0.003416473
Alaska	0.50786248	1.1068225	-1.2117642	2.484202941
Arizona	0.07163341	1.4788032	0.9989801	1.042878388
Arkansas	0.23234938	0.2308680	-1.0735927	-0.184916602
California	0.27826823	1.2628144	1.7589234	2.067820292
Colorado	0.02571456	0.3988593	0.8608085	1.864967207

> #Build kmeans model

> km <- kmeans(data_standardized, centers=2, nstart=20)

> cluster_df <- data_standardized

> cluster_df\$Cluster <- km\$cluster

> cluster_df

	Murder	Assault	UrbanPop	Rape	Cluster
--	--------	---------	----------	------	---------

Alabama	1.24256408	0.78283935	-0.52090661	-0.003416473	1
Alaska	0.50786248	1.10682252	-1.21176419	2.484202941	1
Arizona	0.07163341	1.47880321	0.99898006	1.042878388	1
Arkansas	0.23234938	0.23086801	-1.07359268	-0.184916602	2
California	0.27826823	1.26281442	1.75892340	2.067820292	1
Colorado	0.02571456	0.39885929	0.86080854	1.864967207	1
Connecticut	-1.03041900	-0.72908214	0.79172279	-1.081740768	2
Delaware	-0.43347395	0.80683810	0.44629400	-0.579946294	2
Florida	1.74767144	1.97077766	0.99898006	1.138966691	1
Georgia	2.20685994	0.48285493	-0.38273510	0.487701523	1
Hawaii	-0.57123050	-1.49704226	1.20623733	-0.110181255	2
Idaho	-1.19113497	-0.60908837	-0.79724965	-0.750769945	2
Illinois	0.59970018	0.93883125	1.20623733	0.295524916	1
Indiana	-0.13500142	-0.69308401	-0.03730631	-0.024769429	2
Iowa	-1.28297267	-1.37704849	-0.58999237	-1.060387812	2
Kansas	-0.41051452	-0.66908525	0.03177945	-0.345063775	2
Kentucky	0.43898421	-0.74108152	-0.93542116	-0.526563903	2
Louisiana	1.74767144	0.93883125	0.03177945	0.103348309	1
Maine	-1.30593210	-1.05306531	-1.00450692	-1.434064548	2
Maryland	0.80633501	1.55079947	0.10086521	0.701231086	1
Massachusetts	-0.77786532	-0.26110644	1.34440885	-0.526563903	2
Michigan	0.99001041	1.01082751	0.58446551	1.480613993	1
Minnesota	-1.16817555	-1.18505846	0.03177945	-0.676034598	2
Mississippi	1.90838741	1.05882502	-1.48810723	-0.441152078	1
Missouri	0.27826823	0.08687549	0.30812248	0.743936999	1
Montana	-0.41051452	-0.74108152	-0.86633540	-0.515887425	2
Nebraska	-0.80082475	-0.82507715	-0.24456358	-0.505210947	2
Nevada	1.01296983	0.97482938	1.06806582	2.644350114	1
New Hampshire	-1.30593210	-1.36504911	-0.65907813	-1.252564419	2
New Jersey	-0.08908257	-0.14111267	1.62075188	-0.259651949	2
New Mexico	0.82929443	1.37080881	0.30812248	1.160319648	1
New York	0.76041616	0.99882813	1.41349461	0.519730957	1
North Carolina	1.19664523	1.99477641	-1.41902147	-0.547916860	1
North Dakota	-1.60440462	-1.50904164	-1.48810723	-1.487446939	2
Ohio	-0.11204199	-0.60908837	0.65355127	0.017936483	2
Oklahoma	-0.27275797	-0.23710769	0.16995096	-0.131534211	2
Oregon	-0.66306820	-0.14111267	0.10086521	0.861378259	2
Pennsylvania	-0.34163624	-0.77707965	0.44629400	-0.676034598	2
Rhode Island	-1.00745957	0.03887798	1.48258036	-1.380682157	2
South Carolina	1.51807718	1.29881255	-1.21176419	0.135377743	1
South Dakota	-0.91562187	-1.01706718	-1.41902147	-0.900240639	2

Tennessee	1.24256408	0.20686926	-0.45182086	0.605142783	1
Texas	1.12776696	0.36286116	0.99898006	0.455672088	1
Utah	-1.05337842	-0.60908837	0.99898006	0.178083656	2
Vermont	-1.28297267	-1.47304350	-2.31713632	-1.071064290	2
Virginia	0.16347111	-0.17711080	-0.17547783	-0.056798864	2
Washington	-0.86970302	-0.30910395	0.51537975	0.530407436	2
West Virginia	-0.47939280	-1.07706407	-1.83353601	-1.273917376	2
Wisconsin	-1.19113497	-1.41304662	0.03177945	-1.113770203	2
Wyoming	-0.22683912	-0.11711392	-0.38273510	-0.601299251	2

>

> #Q.19 Print the cluster number for each observation and cluster size for the above k-means model.

```
> table(cluster_df$Cluster)
```

1 2

20 30

>

> #Q.20 Plot the result of the k-means cluster.

```
> library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
> fviz_cluster(km, data=cluster_df[, -5], palette=c("aquamarine", "orange"), geom="point",
ellipse.type="convex", ggtheme=theme_bw())
```

